EveryMind 988 AI process appendix

We used Claude and Codex to help build and review the survey analysis pipeline, dashboard, presentation, demo script, and written deliverables. The demo uses a system-generated voice. AI assisted implementation and interpretation; source-data checks and human comment review exposed errors and limited the claims we could support.

How AI was used

AI helped write cleaning and validation code, develop keyword and regular-expression theme rules, implement the offline dashboard, draft recommendations, and review numerical claims. The operational pipeline uses inspectable rules rather than a generative model to classify each new comment. Changes and unresolved review issues are recorded in the project audit files.

Representative prompts and corrections

The prompt “give the validated data a look and make nessacry adjustments with it” directed review of human annotations. Feedback led to narrower outreach rules: “bring more activities” no longer counts as reaching new audiences. Spanish-language matches were also corrected. The instruction “just consider it marked as y” applied only to review entry B-166; only that entry’s reviewed field was changed.

Other corrections preserved the literal response “None” during CSV import, fixed date parsing, and expanded presenter-name redaction after manual review found missed names. Code review also found partial-output publishing risks and incomplete workbook validation; these were corrected. An invalid reviewer code is excluded rather than silently repaired.

A later prompt asked Codex to “look at claudes work that he just pushed and let me know what you think about the dashboard and his conclusions”. That review found that session “predictions” were descriptive benchmarks and that the student recommendation gap did not establish weaker learning. The revised dashboard labels these comparisons exploratory and separates recommendation scores from learning claims.

How outputs were checked

A separate validation script checked cleaned values against the raw survey, including workbook cells; 100 checks passed in the recorded run. Theme and dashboard suites passed 72 and 33 tests, respectively. Browser review checked desktop and mobile filters. These checks support implementation correctness, not the causal validity of recommendations.

Three human reviewers marked 600 assigned entries reviewed, covering 560 distinct comments. One typo leaves 599 usable entries; five intent/theme conflicts are withheld from theme scoring. Theme precision is 99.2% and recall 98.6% across 594 entries; intent agreement is 294/300 (98.0%). Reviewers saw proposed codes, and the same comments informed rule revisions. These figures describe calibration agreement, not blind or held-out accuracy.

Remaining limits

Reviewer disagreements remain documented. Recalled before/after ratings measure perceived change; missing answers and possible prompt differences limit comparisons. Proposed pilots, budgets, and targets require staff confirmation. Raw survey data and identifying files remain outside the public repository; the authorized validation samples are an exception.